


# Towards Soft Error Resilience in SWE with TeaMPI

## Bachelor's Thesis Talk

**Atamert Rahma**

Chair of Scientific Computing in Computer Science  
Department of Informatics  
Technical University of Munich

September 29<sup>th</sup>, 2021



- 1 Failures in HPC
- 2 SWE with TeaMPI
- 3 Integrating Soft Error Resilience Techniques into SWE
- 4 Comparison of Different Soft Error Resilience Techniques
- 5 Conclusion

# Failures in HPC

**Hard Errors**  $\Rightarrow$  stop execution

**Soft Errors**  $\Rightarrow$  no permanent failure

# Failures in HPC

**Hard Errors**  $\Rightarrow$  stop execution

**Soft Errors**  $\Rightarrow$  no permanent failure

Cosmic radiation  $\rightarrow$  Energetic particles (Neutrons, Alpha particles) hit the silicon device  $\rightarrow$   
Cause a sufficient charge  $\rightarrow$  Inverts the state of a logic device (bitflip)  $\rightarrow$  Soft error

Soft errors may lead to

- **DUE** (Detectable and Uncorrectable Error)
- **SDC** (Silent Data Corruption)

# Failures in HPC

**Hard Errors**  $\Rightarrow$  stop execution

**Soft Errors**  $\Rightarrow$  no permanent failure

Cosmic radiation  $\rightarrow$  Energetic particles (Neutrons, Alpha particles) hit the silicon device  $\rightarrow$   
Cause a sufficient charge  $\rightarrow$  Inverts the state of a logic device (bitflip)  $\rightarrow$  Soft error

Soft errors may lead to

- **DUE** (Detectable and Uncorrectable Error)
- **SDC** (Silent Data Corruption)

Increased number of system components (CPU, memory)  $\Rightarrow$  **higher error rates!**

# Fault Tolerance in HPC

1. **CR** (Checkpoint/Restart)
2. **Replication** (of threads or processes)

# Fault Tolerance in HPC

1. **CR** (Checkpoint/Restart)
2. **Replication** (of threads or processes)

CR overhead is increasing  $\Rightarrow$  CR is not expected to be the best solution in the future

Checkpoints may include SDCs  $\Rightarrow$  **CR alone cannot provide soft error resilience!**

# Outline

- 1 Failures in HPC
- 2 SWE with TeaMPI**
- 3 Integrating Soft Error Resilience Techniques into SWE
- 4 Comparison of Different Soft Error Resilience Techniques
- 5 Conclusion




# SWE for Solving Shallow Water Equations

SWE is a teaching code that can simulate

- different wave propagation/tsunami scenarios
- supports different parallel processing models (**MPI**, CUDA)

⇒ An example scenario where 4 MPI ranks were used.

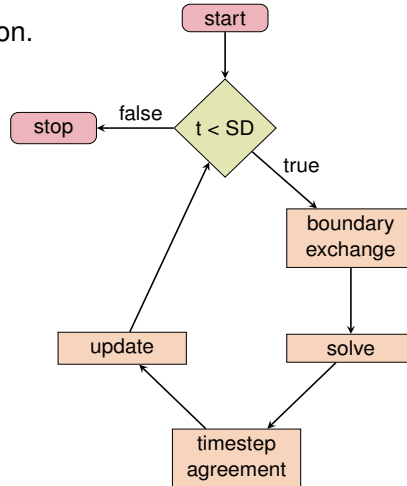


# SWE for Solving Shallow Water Equations

Computation loop of a simple SWE application.

■ **t** = time

■ **SD** = simulation duration



Wrapper library for MPI that utilizes **process replication**.

Wrapper library for MPI that utilizes **process replication**.

- **Heartbeats**: non-blocking messages between the replicated ranks
- **Task Sharing**: to reduce redundant computation time

- 1 Failures in HPC
- 2 SWE with TeaMPI
- 3 Integrating Soft Error Resilience Techniques into SWE**
- 4 Comparison of Different Soft Error Resilience Techniques
- 5 Conclusion

# Soft Error Detection with Hashes

## *Hashes*

Idea: **Process replication** + **Hash value comparison** of the redundantly computed results

# Soft Error Detection with Hashes

## *Hashes*

Idea: **Process replication** + **Hash value comparison** of the redundantly computed results

- Data arrays of our application: **bathymetry**, **water height** and **momentum arrays**

# Soft Error Detection with Hashes

## *Hashes*

Idea: **Process replication** + **Hash value comparison** of the redundantly computed results

- Data arrays of our application: **bathymetry**, **water height** and **momentum arrays**
- **Single heartbeats** carry the hash values to replicas



# Soft Error Detection with Hashes

## *Hashes*

Idea: **Process replication** + **Hash value comparison** of the redundantly computed results

- Data arrays of our application: **bathymetry**, **water height** and **momentum arrays**
- **Single heartbeats** carry the hash values to replicas
- Comparison is handled **transparently** in TeaMPI

# Soft Error Detection with Hashes

## *Hashes*

Idea: **Process replication** + **Hash value comparison** of the redundantly computed results

- Data arrays of our application: **bathymetry**, **water height** and **momentum arrays**
- **Single heartbeats** carry the hash values to replicas
- Comparison is handled **transparently** in TeaMPI
- A mismatching hash value is assumed to be a sign of an SDC

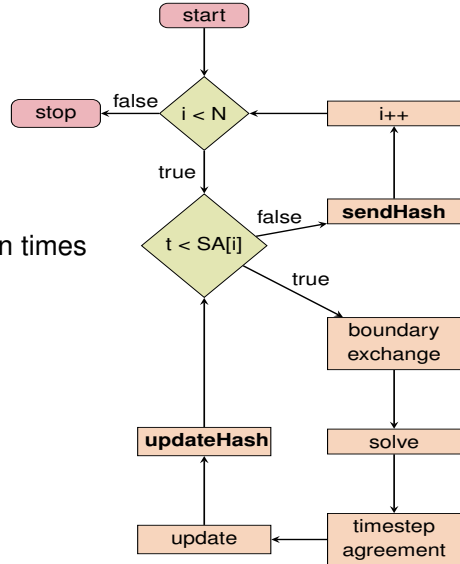
⇒ can provide SDC detection!

# Hashes – Implementation

Computation loop for *Hashes*

- **t** = time, **N** = total hash sends
- **SA** = SendAt = equally distanced simulation times to send the hashes
- **i** = counter

This can only provide soft error detection!



# Soft Error Resilience Using Verification and Task Sharing

## *Sharing*



Idea: **Process replication + Admissibility checks + Task sharing**

# Soft Error Resilience Using Verification and Task Sharing

## *Sharing*



Idea: **Process replication** + **Admissibility checks** + **Task sharing**

- Admissibility validation of the data and **update** arrays in our solver scheme

# Soft Error Resilience Using Verification and Task Sharing

## *Sharing*



Idea: **Process replication** + **Admissibility checks** + **Task sharing**

- Admissibility validation of the data and **update** arrays in our solver scheme
- Recovery using one of the "healthy" replicas in case of violation

# Soft Error Resilience Using Verification and Task Sharing

## *Sharing*



Idea: **Process replication** + **Admissibility checks** + **Task sharing**

- Admissibility validation of the data and **update** arrays in our solver scheme
- Recovery using one of the "healthy" replicas in case of violation
- **Primary** and **secondary** blocks

# Soft Error Resilience Using Verification and Task Sharing

## *Sharing*



Idea: **Process replication** + **Admissibility checks** + **Task sharing**

- Admissibility validation of the data and **update** arrays in our solver scheme
- Recovery using one of the "healthy" replicas in case of violation
- **Primary** and **secondary** blocks

⇒ An undetected SDC may **propagate** to the replicated ranks and corrupt them as well due to task sharing!



# Admissibility Criteria

1. Physical Admissibility Criteria
  - ☐ Constant Bathymetry
  - ☐ Non-negative Water Height
2. Numerical Admissibility Criteria
  - ☐ No Float Errors (NaNs)

# Admissibility Criteria

## 1. Physical Admissibility Criteria

- ☐ Constant Bathymetry
- ☐ Non-negative Water Height

## 2. Numerical Admissibility Criteria

- ☐ No Float Errors (NaNs)
- ☐ Relaxed **DMP** (Discrete Maximum Principle)

$$\underbrace{\min_{y \in V_{i,j}} u(y, t^n) - \delta}_{\text{minimum neighbor}} \leq \overbrace{u^*(x, t^{n+1})}^{\text{candidate solution}} \leq \underbrace{\max_{y \in V_{i,j}} u(y, t^n) + \delta}_{\text{maximum neighbour}}$$

# Admissibility Criteria

## 1. Physical Admissibility Criteria

- ☐ Constant Bathymetry
- ☐ Non-negative Water Height

## 2. Numerical Admissibility Criteria

- ☐ No Float Errors (NaNs)
- ☐ Relaxed **DMP** (Discrete Maximum Principle)

$$\underbrace{\min_{y \in V_{i,j}} u(y, t^n) - \delta}_{\text{minimum neighbor}} \leq \overbrace{u^*(x, t^{n+1})}^{\text{candidate solution}} \leq \underbrace{\max_{y \in V_{i,j}} u(y, t^n) + \delta}_{\text{maximum neighbour}}$$

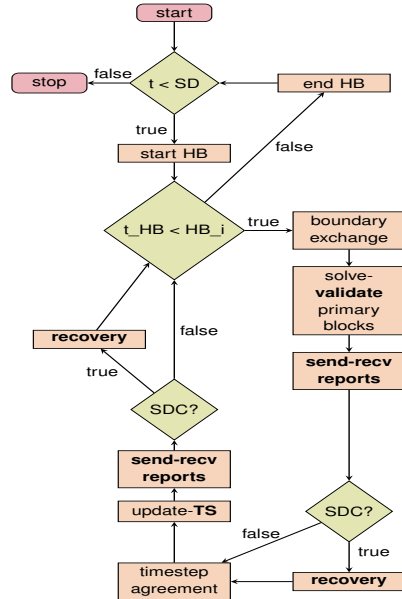
⇒ Relaxed DMP is not a strong condition and can give false positives!

# Sharing – Implementation

Computation loop for *Sharing*

- $t$  = time,  $SD$  = simulation duration
- $t_{HB}$  = time since last heartbeat
- $HB_i$  = heartbeat interval
- $TS$  = task sharing

Additional soft error resilience with recovery!



# Soft Error Resilience Using Redundant Computation

## *Redundant*



Idea: **Process replication** + **Admissibility checks**

- Only **primary** blocks
- Teams run independently (**no task sharing**)

# Soft Error Resilience Using Redundant Computation

## *Redundant*

Idea: **Process replication** + **Admissibility checks**

- Only **primary** blocks
- Teams run independently (**no task sharing**)

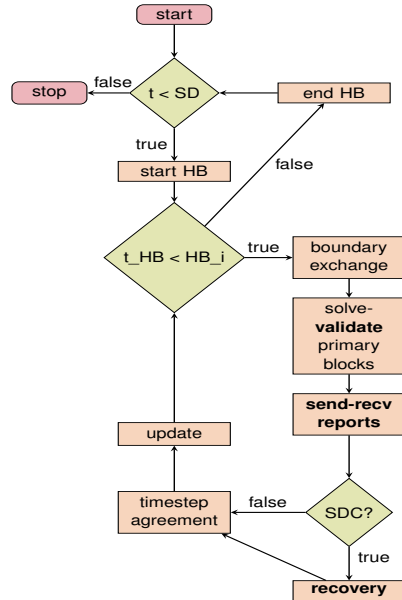
⇒ A possible SDC in one team **cannot propagate** to other teams!

# Redundant – Implementation

Computation loop for *Redundant*

- $t$  = time,  $SD$  = simulation duration
- $t_{HB}$  = time since last heartbeat
- $HB_i$  = heartbeat interval

Additional soft error resilience with redundant computation!



# Outline

- 1 Failures in HPC
- 2 SWE with TeaMPI
- 3 Integrating Soft Error Resilience Techniques into SWE
- 4 Comparison of Different Soft Error Resilience Techniques**
- 5 Conclusion



# Bitflip Injections

We inject only a single bitflip per run. Bitflips are injected

- at a fixed location in the source code,
- into the data and update arrays at a random number and bit.

# Soft Error Outcome Rates

## Random Bitflip Injections

10000 injections into each listed array below (random injection with 30000 injections).

*Sharing*  $\Rightarrow$

<i>outcome\array</i>	<i>h</i>	<i>hu</i>	<i>hv</i>	updates	random
Correctable	18.55%	13.83%	14.47%	<b>0%</b>	15.22%
DUE	0.25%	0.13%	0.04%	7.32%	4.36%
SDC	81.21%	86.03%	85.49%	92.68%	80.41%
Negligible	2.95%	3.26%	3.31%	28.17%	19.84%

*Redundant*  $\Rightarrow$

<i>outcome\array</i>	<i>h</i>	<i>hu</i>	<i>hv</i>	updates	random
Correctable	18.6%	13.63%	15%	<b>6.09%</b>	19.06%
DUE	0.03%	0.1%	0.04%	0.96%	0.43%
SDC	81.37%	86.26%	84.96%	92.95%	80.52%
Negligible	2.91%	3.32%	2.95%	28.09%	19.56%

# Soft Error Outcome Rates

## Relaxation Factor ( $\delta$ )

Correctable outcome rates of 2000 injections into the leftmost 10 bits of randomly selected floats.

$\delta \backslash array$	<i>Sharing</i>				<i>Redundant</i>			
	<i>h</i>	<i>hu</i>	<i>hv</i>	updates	<i>h</i>	<i>hu</i>	<i>hv</i>	updates
80	<b>59.44%</b>	<b>43.37%</b>	<b>50.44%</b>	<b>0%</b>	<b>60.34%</b>	<b>45.31%</b>	<b>50.7%</b>	<b>15.76%</b>
100	59.95%	43.47%	47.07%	0%	58.97%	43.91%	47.74%	15.34%
10000	39.69%	29.91%	29.24%	0%	39.96%	28.1%	30.62%	11.21%
1000000	<b>34.9%</b>	<b>22.07%</b>	<b>25.32%</b>	<b>0%</b>	<b>37.68%</b>	<b>24%</b>	<b>25.46%</b>	<b>10.16%</b>

# Soft Error Outcome Rates

## Type of the Simulation Scenario (s)


Correctable outcome rates of 2000 injections into the leftmost 10 bits of randomly selected floats.

- A = *radialBathymetryDamBreak* → lower water height (~15.1 meters)
- B = *splashingPool* → higher water height (~245 meters)
- C = *seaAtRest* → constant water height (10 meters)

Sharing					Redundant			
<i>s\array</i>	<i>h</i>	<i>hu</i>	<i>hv</i>	updates	<i>h</i>	<i>hu</i>	<i>hv</i>	updates
A	59.95%	43.47%	47.07%	0%	58.97%	43.91%	47.74%	15.34%
B	89.77%	89.39%	90.28%	0%	99.95%	90.1%	90.5%	15.99%
C	58.25%	0%	0%	0%	61.65%	0%	0%	0%
C (r=0)	100%	100%	100%	0%	100%	100%	100%	99.8%



# Performance Comparison – Single Node

Profiling on a single node. (**noRes** = no resilience)



# Performance Comparison – Multiple Nodes


Teams are mapped onto the same node vs. Replicas are mapped onto the same node



⇒ increased communication between teams due to task sharing!

# Performance Comparison – Multiple Nodes

## Strong Scaling – replicas on the same node



# Outline

- 1 Failures in HPC
- 2 SWE with TeaMPI
- 3 Integrating Soft Error Resilience Techniques into SWE
- 4 Comparison of Different Soft Error Resilience Techniques
- 5 Conclusion**



## Summary & Future Work

1. *Hashes*: Soft error **detection** using hash comparison
2. *Sharing*: Soft error recovery using **validation** and **task sharing**
3. *Redundant*: Soft error recovery using **validation** and **redundant computation**

## Summary & Future Work

1. *Hashes*: Soft error **detection** using hash comparison
2. *Sharing*: Soft error recovery using **validation** and **task sharing**
3. *Redundant*: Soft error recovery using **validation** and **redundant computation**

⇒ Resilience depends on different factors ( $\delta$ , simulation duration and type).

⇒ *Redundant* can **additionally recover from some SDC cases** that *Sharing* cannot.

⇒ *Sharing* has the **the best performance**.

## Summary & Future Work

1. *Hashes*: Soft error **detection** using hash comparison
2. *Sharing*: Soft error recovery using **validation** and **task sharing**
3. *Redundant*: Soft error recovery using **validation** and **redundant computation**

⇒ Resilience depends on different factors ( $\delta$ , simulation duration and type).  
⇒ *Redundant* can **additionally recover from some SDC cases** that *Sharing* cannot.  
⇒ *Sharing* has the **the best performance**.

Future work:

- hard error resilience evaluation together with our resilience techniques
- more random bitflip injections (different areas in the code)
- additional admissibility criteria and improved recovery to cover more SDC cases